

# HW3 Writeup

## M1522.001000 Computer Vision

Junyeong Park (2017-19460)

---

### Part 1

#### Part 1.1 $\mathbf{H} = \text{compute\_h}(p1, p2)$

```
1 def compute_h(p1, p2):
2     # TODO ...
3     """
4     svd(A) : Gets U, S, V^T that satisfies A = U S V^T
5     S : Eigenvalue of A^T*A
6     The eigenvector when the eigenvalue is 0 can be calculated by V_T[-1, :]
7     """
8     n = len(p1) - 1
9     A = [] # 2N x 9 matrix
10    # Build A
11    while n > 0:
12        A.append([p2[n][0], p2[n][1], 1, 0, 0, 0, -p1[n][0] * p2[n][0], -p1[n][0] * p2[n][1], -p1[n]
13        ↪ ] [0]])
14        A.append([0, 0, 0, p2[n][0], p2[n][1], 1, -p1[n][1] * p2[n][0], -p1[n][1] * p2[n][1], -p1[n]
15        ↪ ] [1]])
16        n -= 1
17    A = np.asarray(A)
18
19    # SVD
20    U, S, V_T = np.linalg.svd(A)
21    h = V_T[-1, :]
22    H = np.reshape(h, (3, 3))
23    # Normalize
24    H /= H[2][2]
25
26    return H
```

Line 8~15 are for building  $\mathbf{A}$ , a  $2N \times 9$  matrix. Line 17~24 are for building  $\mathbf{H}$ , a homography matrix. The expressions of these matrices are mentioned on **Theory Question 3**.

To get  $\mathbf{H}$ , we have to find  $\mathbf{h}$  that minimizes  $|\mathbf{A}\mathbf{t}|^2 = (\mathbf{A}\mathbf{h})^T(\mathbf{A}\mathbf{h}) = \lambda$ . After we take a derivative w.r.t.  $\mathbf{h}$  on this equation, we get a result that  $\lambda$  is an eigenvalue of  $\mathbf{A}^T\mathbf{A}$ , and the solution  $\mathbf{h}$  is its eigenvector.

Getting  $\mathbf{h}$  that minimizes  $\lambda$  can be done easily by using SVD, which stands for singular value decomposition. `np.linalg.svd(A)` returns three matrices  $\mathbf{U}$ ,  $\mathbf{S}$ ,  $\mathbf{V}^T$ .  $\mathbf{S}$  is a diagonal matrix that contains eigenvalues of  $\mathbf{A}^T\mathbf{A}$ , and  $\mathbf{V}$  is a matrix that contains the corresponding eigenvectors. We can get  $\mathbf{h}$  that makes  $\lambda = 0$  from the last row of  $\mathbf{V}^T$ . After reshaping  $\mathbf{h}$ , we get an unnormalized homography matrix. By dividing  $\mathbf{H}$  by  $h_{22}$ , we finally get a normalized homography matrix  $\mathbf{H}$ .

#### Part 1.2 $\mathbf{H} = \text{compute\_h\_norm}(p1, p2)$

```
def compute_h_norm(p1, p2):
2     # TODO ...
3     # Transformation matrix
4     # (x, y) -> (x/col, y/row)
5     T = [[1/400, 0, 0], [0, 1/302, 0], [0, 0, 1]]
6     T = np.asarray(T)
```

```

8     inv_T = np.linalg.inv(T)
10
11     p1_norm, p2_norm = p1, p2
12     # Compute normalized coordinates
13     for i in range(len(p1)):
14         p1_coord = np.matmul(T, [p1_norm[i][0], p1_norm[i][1], 1])
15         p2_coord = np.matmul(T, [p2_norm[i][0], p2_norm[i][1], 1])
16         p1_norm[i][0], p1_norm[i][1] = p1_coord[0], p1_coord[1]
17         p2_norm[i][0], p2_norm[i][1] = p2_coord[0], p2_coord[1]
18
19     H_norm = compute_h(p1_norm, p2_norm)
20
21     # H = T^-1 * H_norm * T
22     H = np.matmul(np.matmul(inv_T, H_norm), T)
23     # Floating point problem
24     H /= H[2][2]
25
26     return H

```

The basic idea of this function is actually the same with `compute_h`. The difference is that there is a transformation matrix  $\mathbf{T}$  in `compute_h_norm`.  $\mathbf{T}$  is a matrix that normalizes the coordinates by scaling  $x$  by  $\frac{1}{400}$  and  $y$  by  $\frac{1}{302}$ . 400 and 302 are the width and height of the images, i.e. `wdc1.png` and `wdc2.png`, respectively. After computing  $\mathbf{H}_{\text{norm}}$  from two normalized coordinates, we can finally get  $\mathbf{H} = \mathbf{T}^{-1}\mathbf{H}_{\text{norm}}\mathbf{T}$ .

Since it is not able to infer the shape of `igs_in` and `igs_ref` from the arguments, I had to hard-code parameters of the width and the height of the images.

## Part 2

### Part 2.1 `p_in, p_ref = set_cor_mosaic()`

```

1 def set_cor_mosaic():
2     # TODO ...
3     """
4     p_in : N x 2 matrices of corresponded (x, y)^T coordinates in igs_in
5     p_ref : N x 2 matrices of corresponded (x, y)^T coordinates in igs_ref
6     N=13 point correspondences are chosen.
7     Points that I chose are mentioned in the writeup.
8     """
9     p_in = [
10         [370, 244],
11         [155, 162],
12         [214, 190],
13         [164, 125],
14         [339, 92],
15         [173, 187],
16         [217, 161],
17         [88, 174],
18         [223, 103],
19         [384, 137],
20         [359, 194],
21         [357, 110],
22         [325, 288]
23     ]
24
25     p_ref = [
26         [111, 98],
27         [347, 68],
28         [272, 74],
29         [387, 97],
30         [259, 253],
31         [305, 58],
32         [298, 90],

```

```

34         [376, 34],
35         [368, 149],
36         [161, 209],
37         [146, 135],
38         [218, 232],
39         [130, 58]
40     ]
41
42     return p_in, p_ref

```

I chose 13 point correspondences from the images. Figure 1 denotes the points that I chose. Since the function doesn't get anything as an argument, I decided to hard-code these points. Theoretically, 4 point correspondences are enough. However, the more correspondences lead to a better result, because we use the method of least squares to compute  $\mathbf{H}$ . Interestingly, I found out that choosing points sparsely also leads to a better result.




Figure 1: 13 point correspondences on wdc2.png

## Part 2.2 `igs_warp, igs_merge = warp_image(igs_n, igs_ref, H)`

```

1 def warp_image(igs_in, igs_ref, H):
2     # TODO ...
3     inv_H = np.linalg.inv(H) # Inverse warping
4     in_y, in_x, _ = igs_in.shape
5     ref_y, ref_x, _ = igs_ref.shape
6
7     # Corners of igs_in & igs_ref
8     in_corners = [(0, 0), (in_y, in_x), (0, in_x), (in_y, 0)]
9     ref_corners = [(0, 0), (ref_y, ref_x), (0, ref_x), (ref_y, 0)]
10
11    # Maximum and minimum values of x & y in igs_warp
12    min_x, min_y = float("-inf"), float("-inf")
13    max_x, max_y = float("inf"), float("inf")
14
15    # Compute min, max
16    for i, j in in_corners:
17        # Homogeneous coordinate
18        x, y, w = np.matmul(H, [j, i, 1])
19        x = x/w
20        y = y/w

```

```

21     if x > max_x:
22         max_x = int(x)
23     if x < min_x:
24         min_x = int(x)
25     if y > max_y:
26         max_y = int(y)
27     if y < min_y:
28         min_y = int(y)
29
30     igs_warp = np.zeros((max_y - min_y, max_x - min_x, 3))
31
32     # Compute igs_warp
33     for i in range(0, max_x - min_x):
34         for j in range(0, max_y - min_y):
35             # Homogeneous coordinate
36             x, y, w = np.matmul(inv_H, [i + min_x, j + min_y, 1])
37             # interpolate
38             x = int(x/w)
39             y = int(y/w)
40             # Colors
41             r, g, b = 0, 0, 0
42             if not (y < 0 or y >= in_y or x < 0 or x >= in_x):
43                 r, g, b = igs_in[y, x, :]
44             igs_warp[j, i, :] = [r, g, b]
45
46     # old : min & max of igs_warp
47     old_min_x = min_x
48     old_min_y = min_y
49     old_max_x = max_x
50     old_max_y = max_y
51
52     # Compute min, max of igs_merge
53     for i, j in ref_corners:
54         if j > max_x:
55             max_x = int(j)
56         if j < min_x:
57             min_x = int(j)
58         if i > max_y:
59             max_y = int(i)
60         if i < min_y:
61             min_y = int(i)
62
63     igs_merge = np.zeros(((max_y - min_y), (max_x - min_x), 3))
64
65     # Compute igs_merge
66     for i in range(min_x, max_x):
67         for j in range(min_y, max_y):
68             r, g, b = 0, 0, 0
69             # igs_warp
70             if not (j < old_min_y or j >= old_max_y or i < old_min_x or i >= old_max_x):
71                 r, g, b = igs_warp[j - old_min_y, i - old_min_x, :]
72                 if r * g * b == 0.0:
73                     if not (j < 0 or j >= ref_y or i < 0 or i >= ref_x):
74                         r, g, b = igs_ref[j, i, :]
75             # igs_ref
76             else:
77                 if not (j < 0 or j >= ref_y or i < 0 or i >= ref_x):
78                     r, g, b = igs_ref[j, i, :]
79             igs_merge[j - min_y, i - min_x, :] = [r, g, b]
80
81     return igs_warp, igs_merge
82
83

```

In this function, we get `igs_warp`, a warped version of `igs_in`, and `igs_merge`, a single mosaic image with a larger field of view containing both of the input images. Figure 2 is a mosaic image generated from `igs_merge`. I used an inverse warping technique, and interpolated the non-integer valued coordinates by type-casting them as `int()`. I tried to use other interpolation techniques such as bilinear interpolation, but I decided not to use them

because they didn't really helped the quality of the output image and also took a lot of computation time to get the result. The rest of the code is explained in the comments.




Figure 2: Mosaic image

## Part 3

### Part 3.1 `c_in, c_ref = set_cor_rec()`

```

def set_cor_rec():
2     # TODO ...
    """
4     p_in : N x 2 matrices of corresponded (x, y)^T coordinates in igs_in
    p_ref : N x 2 matrices of corresponded (x, y)^T coordinates in igs_ref
6     N=4 point correspondences are chosen.
    Points that I chose are mentioned in the writeup.
    """
8
    c_in = [
10         [163, 14],
11         [260, 25],
12         [163, 258],
13         [260, 245]
14     ]

    c_ref = [
16         [200, 70],
17         [260, 1],
18         [150, 252],
19         [258, 250]
20     ]

22     return c_in, c_ref
24

```

Figure 3 denotes the points that I chose. The points where red lines meet consist `c_in`. For `c_out`, I adjusted the points manually to get a better result.

### Part 3.2 `igs_rec = rectify(igs, p1, p2)`



Figure 3: iphone.png



Figure 4: Rectified version of Figure 3

```

1 def rectify(igs, p1, p2):
    # TODO ...
    # Get homography matrix of p1->p2
    H = compute_h(p2, p1)
    # Use warp_image to get igs_rec
    igs_rec, _ = warp_image(igs, np.zeros(igs.shape), H)
    7
    return igs_rec
    9

```

In this function, I reused the functions from **Part 1** and **Part 2**. Figure 4 is an image generated from `igs_rec`. The problem is that I was not able to use `compute_h_norm` because the transformation matrix fits exclusively for the images in **Part 2**. This led to the manual adjustment on `c_ref`, since the scale is different on the input and the output.